

1 Equivalence of Matrix Norm Definitions

(The purpose of this question is to improve your comfort in manipulating matrix algebra.)

The squared Frobenious norm has two equivalent definitions:

$$\|X\|_F^2 = \text{trace}(X^T X), \quad \text{and} \quad \|X\|_F^2 = \sum_d \sigma_d^2, \quad (1)$$

where σ_d denotes the d -th singular value of X , i.e., σ_d forms the diagonal of Σ in the SVD of $X = U\Sigma V^T$.

Similarly, the trace norm has two equivalent definitions:

$$\|X\|_* = \text{trace}(\sqrt{X^T X}), \quad \text{and} \quad \|X\|_* = \sum_d \sigma_d, \quad (2)$$

where for any symmetric square matrix A , if $B = \sqrt{A}$, then $BB = A$.

Question: Prove that the two expressions in (1) are equivalent, and prove that the two expressions in (2) are equivalent. You can assume that X is a square matrix for convenience.

Solution. We will leverage the SVD of $X = U\Sigma V^T$. To show (1), we see that

$$\text{trace}(X^T X) = \text{trace}(V\Sigma U^T U\Sigma V^T) \quad (3)$$

$$= \text{trace}(V\Sigma V^T) \quad (4)$$

$$= \text{trace}(V\Sigma^2 V^T) \quad (5)$$

$$= \text{trace}(\Sigma^2 V^T V) \quad (6)$$

$$= \text{trace}(\Sigma^2) \quad (7)$$

$$= \sum_d \sigma_d^2, \quad (8)$$

where (1) follows from the fact that U is orthogonal and thus $U^T U = I$, (6) follows from the rotational invariance of the trace operator: $\text{trace}(ABC) = \text{trace}(BCA) = \text{trace}(CAB)$, and (7) follows from the fact that V is orthogonal and thus $V^T V = I$.

To show (2), we see that

$$\text{trace}(\sqrt{X^T X}) = \text{trace}(\sqrt{V\Sigma U^T U\Sigma V^T}) \quad (9)$$

$$= \text{trace}(\sqrt{\Sigma^2}) \quad (10)$$

$$= \text{trace}(\Sigma) \quad (11)$$

$$= \sum_d \sigma_d, \quad (12)$$

where (10) follows from the same logic as (4)-(7).

2 Properties of Bootstrap Sampling

(The purpose of this question is to improve your comfort in manipulating probability concepts.)

Let $S = \{(x_i, y_i)\}_i = 1^N$ denote a training set of size N . In this question, we will analyze the properties of a single bootstrapped dataset of S . Recall that Bootstrapping is the process of creating a new dataset S' of size N where each entry in S' is sampled uniformly from S (with replacement). Procedurally, this means

- Initialize $S' \leftarrow \emptyset$
- For $j = 1, \dots, N$
 - sample (x'_j, y'_j) by sampling uniformly from S (and independently of other samples)
 - add (x'_j, y'_j) to S'
- Return S'

This means that some entries in S will appear multiple times in S' , and some entries won't appear at all. For a specific $(x, y) \in S$, what is the probability that (x, y) will appear at least once in S' .

Question: What does this probability converge to in the limit as N increases? You can assume that every $(x, y) \in S$ is unique (there are no duplicates in S).

Solution. Let $P((x, y) \in S')$ denote the probability that (x, y) appears at least once in S' . Then it suffices to compute $P((x, y) \notin S') = 1 - P((x, y) \in S')$. We can write $P((x, y) \notin S')$ as

$$P((x, y) \notin S') = P(\forall (x'_j, y'_j) \in S' : (x, y) \neq (x'_j, y'_j)) \quad (13)$$

$$= \prod_{j=1}^N P((x, y) \neq (x'_j, y'_j)) \quad (14)$$

$$= \prod_{j=1}^N \left(1 - \frac{1}{N}\right) \quad (15)$$

$$= \left(1 - \frac{1}{N}\right)^N, \quad (16)$$

where (14) follows from the independence of sampling each (x'_j, y'_j) , and (15) is simply using the definition of uniform distribution. Thus we can see that the probability of (x, y) appearing at least once in S' is

$$P((x, y) \in S') = 1 - P((x, y) \notin S') = 1 - \left(1 - \frac{1}{N}\right)^N.$$

In the limit as N increases, we can solve:

$$c \equiv \lim_{N \rightarrow \infty} \log \left(\left(1 - \frac{1}{N}\right)^N \right), \quad (17)$$

after which we know that:

$$\lim_{N \rightarrow \infty} 1 - \left(1 - \frac{1}{N}\right)^N = 1 - e^c. \quad (18)$$

We will use L'Hospital's rule¹ to calculate (17):

$$\begin{aligned}\lim_{N \rightarrow \infty} \log \left(\left(1 - \frac{1}{N} \right)^N \right) &= \lim_{N \rightarrow \infty} N \log \left(1 - \frac{1}{N} \right) \\ &= \lim_{N \rightarrow \infty} \frac{\log \left(1 - \frac{1}{N} \right)}{\frac{1}{N}} \\ &= \lim_{N \rightarrow \infty} \frac{\frac{1}{1 - \frac{1}{N}} \cdot \frac{1}{N^2}}{-\frac{1}{N^2}}\end{aligned}\tag{19}$$

$$\begin{aligned}&= \lim_{N \rightarrow \infty} -\frac{1}{1 - \frac{1}{N}} \\ &= -1,\end{aligned}\tag{20}$$

where (19) follows from applying L'Hospital's rule and the chain rule when differentiating $\log(1 - 1/N)$. Thus we see that

$$\lim_{N \rightarrow \infty} P((x, y) \in S') = 1 - \frac{1}{e}.$$

¹http://en.wikipedia.org/wiki/L%27H%C3%B4pital%27s_rule

3 Bias-Variance Decomposition

(The purpose of this question is to improve your comfort with manipulating probability concepts and loss functions.)

Let f_S denote a regression model² whose parameters are trained using training set S . Then if we treat the training set S as a random variable (i.e., sampled from the true test distribution), then f_S is also a random variable.

For any test data point x with true label y , we can write the expected squared loss of f_S as

$$E_S \left[(f_S(x) - y)^2 \right], \quad (21)$$

where the expectation is taken over the randomness of the training set S .

Question #1: Derive the bias-variance decomposition of (21):

$$E_S \left[(f_S(x) - y)^2 \right] = E_S \left[(f_S(x) - \bar{f}(x))^2 \right] + E_S \left[(\bar{f}(x) - y)^2 \right], \quad (22)$$

where \bar{f} denotes the expectation of f_S over the randomness of S :

$$\bar{f}(x) = E_S [f_S(x)]. \quad (23)$$

Question #2: Why are the two terms in the RHS of (22) referred to as the variance and bias of $f_S(x)$, respectively?

Solution. The first term in the RHS of (22) is called the variance of $f_S(x)$ because it measures the variance of the distribution of $f_S(x)$, and actually does not depend on the true label y at all. The second term in the RHS of (22) is called the bias of $f_S(x)$ because it measures the bias of the distribution of $f_S(x)$ relative to the true label y . In other words, if the bias term were 0, then the distribution of $f_S(x)$ could be interpreted as being “centered” around the true label y .

The derivation is as follows:

$$\begin{aligned} E_S \left[(f_S(x) - y)^2 \right] &= E_S \left[f_S(x)^2 - 2yf_S(x) + y^2 \right] \\ &= E_S \left[f_S(x)^2 \right] - 2yE_S \left[f_S(x) \right] + y^2 \end{aligned} \quad (24)$$

$$= E_S \left[f_S(x)^2 \right] - 2y\bar{f}(x) + y^2 \quad (25)$$

$$= E_S \left[f_S(x)^2 \right] - 2y\bar{f}(x) + y^2 + 2\bar{f}(x)^2 - 2\bar{f}(x)^2 \quad (26)$$

$$= E_S \left[f_S(x)^2 \right] - 2y\bar{f}(x) + y^2 + 2\bar{f}(x) - 2E_S \left[f_S(x) \right] \bar{f}(x) \quad (27)$$

$$= E_S \left[f_S(x)^2 - 2f_S(x)\bar{f}(x) + \bar{f}(x)^2 \right] + E_S \left[\bar{f}(x)^2 - 2y\bar{f}(x) + y^2 \right] \quad (28)$$

$$= E_S \left[(f_S(x) - \bar{f}(x))^2 \right] + E_S \left[(\bar{f}(x) - y)^2 \right].$$

(24) follows from linearity of expectation. (25) follows from applying (23). (26) follows from just adding and subtracting $2\bar{f}(x)^2$ (note that $\bar{f}(x)$ is a deterministic constant given x). (27) follows from applying (23). (28) follows from linearity of expectation and re-arranging terms.

²For any input x , $f_S(x)$ outputs a scalar real value.

4 Multiclass SVM

(The purpose of this question is to improve your comfort in reasoning about **constraints in machine learning optimization problems**.)

Consider the multiclass SVM model that makes predictions on input $x \in \mathbb{R}^D$ via:

$$h(x|w) = \operatorname{argmax}_{y \in \{1, \dots, K\}} w_y^T x,$$

where the model is defined as:

$$w = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_K \end{bmatrix} \in \mathbb{R}^{DK},$$

for each model sub-vector $w_k \in \mathbb{R}^D$.

The multiclass SVM objective over a training set $S = \{(x_i, y_i)\}_{i=1}^N$ can be described as:

$$\operatorname{argmin}_{w, \xi} \frac{1}{2} \|w\|^2 + \frac{C}{N} \sum_{i=1}^N \xi_i \quad (29)$$

s.t.

$$\forall i, \forall y' \in \{1, \dots, K\} : w_{y_i}^T x_i - w_{y'}^T x_i \geq \mathbf{1}_{[y_i \neq y']} - \xi_i \quad (30)$$

Note that $C \geq 0$ is the hyperparameter that trade-off off between regularization and training loss.

Question: For a single training data point (x_i, y_i) , compute the stochastic (sub-)gradient of w by differentiating:

$$\operatorname{argmin}_{w, \xi} \frac{1}{2N} \|w\|^2 + \frac{C}{N} \xi_i \quad (31)$$

s.t.

$$\forall y' \in \{1, \dots, K\} : w_{y_i}^T x_i - w_{y'}^T x_i \geq \mathbf{1}_{[y_i \neq y']} - \xi_i \quad (32)$$

It will be convenient to use the following equivalent definition of ξ_i :

$$\xi_i = \max_{y' \in \{1, \dots, K\}} \{ \mathbf{1}_{[y_i \neq y']} - (w_{y_i}^T x_i - w_{y'}^T x_i) \}. \quad (33)$$

Also, assume for simplicity that exactly one y' is maximal in (33).

Solution. We can write the gradient w.r.t. w of (31) as:

$$\nabla_w = \frac{1}{N} w + \frac{C}{N} \frac{\partial \xi_i}{\partial w}.$$

Let \hat{y} denote the $y' \in \{1, \dots, K\}$ that maximizes the definition in (33) for the current value of w . Then we can simplify the definition of ξ_i as:

$$\xi_i = \mathbf{1}_{[y_i \neq \hat{y}]} - (w_{y_i}^T x_i - w_{\hat{y}}^T x_i), \quad (34)$$

which implies that $\partial \xi_i / \partial w$ is only potentially non-zero in the sub-vectors corresponding to w_{y_i} and $w_{\hat{y}}$.³

In the case where $\hat{y} = y_i$ (i.e., the current model w predicts the correct $\hat{y} = y_i$ with sufficiently large margin), then (34) implies that $\xi_i = 0$ and that

$$\frac{\partial \xi_i}{\partial w} = 0.$$

In the case where $\hat{y} \neq y_i$, then (34) implies that $\xi_i > 0$ and that for each sub-vector w_k :

$$\frac{\partial \xi_i}{\partial w_k} = \begin{cases} x_i & \text{if } k = \hat{y} \\ -x_i & \text{if } k = y_i \\ 0 & \text{otherwise} \end{cases}.$$

Note that this derivation extends to general structured SVMs – it just requires an algorithm to compute \hat{y} since that could take exponential time via exhaustive search.

³For general structured prediction, finding the \hat{y} that maximizes (34) requires us to do something beyond exhaustive enumeration. For instance, for sequence labeling problems \hat{y} can be computed via dynamic programming approaches such as Viterbi.

5 Featurized Latent Factor Models

(The purpose of this question is to improve your comfort in working with more complicated latent-factor models that activate more than one pair of latent factors for each data point.)

Consider a featurized version of collaborative filtering, where we have user features $z \in \mathbb{R}^F$ and item features $x \in \mathbb{R}^D$. We model the rating a user with features z would give to an item with features x as

$$y \approx (Uz)^T(Vx),$$

where $U \in \mathbb{R}^{K \times F}$ and $V \in \mathbb{R}^{K \times D}$ are projection matrices that map user and item features into a K -dimensional latent feature space.⁴

Given training data, our goal is to learn U and V by minimizing the regularized training loss:

$$\operatorname{argmin}_{U,V} \frac{\lambda}{2} (\|U\|_F^2 + \|V\|_F^2) + \frac{1}{2} \sum_{(z,x,y) \in S} (y - (Uz)^T(Vx))^2. \quad (35)$$

Question: Derive the gradient for U .

Solution. The answer is:

$$\partial_U = \lambda U - \sum_{(z,x,y) \in S} [(y - (Uz)^T(Vx)) Vxz^T].$$

The answer can be derived by applying the chain rule to:

$$\frac{\partial}{\partial U} (y - (Uz)^T(Vx))^2 = 2 (y - (Uz)^T(Vx)) \frac{\partial}{\partial U} [-(Uz)^T(Vx)].$$

We next observe that:

$$(Uz)^T(Vx) = \operatorname{trace}(z^T U^T Vx) = \operatorname{trace}(U^T Vxz^T),$$

and use the following matrix derivative identity:

$$\frac{\partial}{\partial U} \operatorname{trace}(U^T M) = M$$

to achieve the result.

⁴This setting reduces to conventional feature-less collaborative filtering by assuming that each z and x have exactly one entry 1 and the rest 0.

6 Tensor Latent Factor Models

(The purpose of this question is to improve your comfort in working with higher order models.)




Figure 1: Illustration of 3-way tensor latent factor model.

Consider a three-way D -dimensional latent factor model as depicted in Figure 1:

$$y_{abc} \approx \langle u_a, v_b, w_c \rangle = \sum_{k=1}^K u_{ak} v_{bk} w_{ck},$$

where $\langle a, b, c \rangle$ denotes a 3-way dot product, and each latent factor u_a, v_b, w_c is a D -dimensional vector that we will estimate from data. Let $S = \{(a, b, c)\}$ denote a set of training indices where we've observed Y_{abc} . Our goal is to learn the model parameters A, B , and C to minimize the regularized training loss over a training set S :

$$\operatorname{argmin}_{U, V, W} \frac{\lambda}{2} (\|U\|_F^2 + \|V\|_F^2 + \|W\|_F^2) + \frac{1}{2} \sum_{(a, b, c) \in S} (Y_{abc} - \langle u_a, v_b, w_c \rangle)^2. \quad (36)$$

where u_a, v_b , and w_c denote the corresponding columns of $U \in \mathbb{R}^{K \times M}$, $V \in \mathbb{R}^{K \times N}$, and $W \in \mathbb{R}^{K \times Q}$, respectively.

Question: Compute the gradient w.r.t. u_a of (36).

Hint: use the Hadamard product⁵ in representing your solution:

$$v_b \circ w_c = \begin{bmatrix} v_{b1} w_{c1} \\ v_{b2} w_{c2} \\ \vdots \\ v_{bK} w_{cK} \end{bmatrix} \in \mathbb{R}^K.$$

Solution. Define $S_a = \{(a', b', c') \in S : a' = a\}$ as the training indices with first index being a . Focusing just on the u_a component, we can rewrite (36) as

$$\operatorname{argmin}_{u_a} \frac{\lambda}{2} \|u_a\|^2 + \frac{1}{2} \sum_{(a, b, c) \in S_a} (Y_{abc} - u_a^T (v_b \circ w_c))^2. \quad (37)$$

⁵http://en.wikipedia.org/wiki/Hadamard_product_%28matrices%29

We can write the gradient of (37) w.r.t. u_a as

$$\partial_{u_a} = \lambda u_a - \sum_{(a,b,c) \in S_a} (Y_{abc} - u_a^T (v_b \circ w_c)) (v_b \circ w_c).$$

Note that this is basically identical to the 2-way latent factor model except for the additional Hadamard product. In other words, when optimizing u_a , we treat $(v_b \circ w_c)$ as the “features” and u_a as the linear model parameter.

7 Neural Net Backprop Gradient Derivation

(The purpose of this question is to improve your comfort with models that have multiple layers.)




Figure 2: Illustration of Simple Neural Network.

In this question, we will consider the following neural network depicted in Figure 2. There are six parameters in the model, $u_1, u_2, w_{11}, w_{21}, w_{12}$, and w_{22} .

The network takes in a 2-dimensional input x and outputs a real value $f(x) \in [0, 1]$. The final output $f(x)$ is a weighted combination of two hidden node activations with a sigmoid transfer function:

$$f(x) = \sigma(u_1 h_1(x) + u_2 h_2(x)), \quad (38)$$

where:

$$\sigma(s) = \frac{e^s}{1 + e^s}.$$

The two hidden layers outputs, $h_1(x)$ and $h_2(x)$, are weighted combinations of the input x with a rectilinear transfer function:

$$h_i(x) = \tau(w_{1i}x_1 + w_{2i}x_2), \quad (39)$$

where:

$$\tau(s) = \max\{0, s\}.$$

Question: For a given training data point (x, y) , compute the stochastic gradient of the squared-loss of (x, y) w.r.t. w_{11} :

$$\frac{\partial}{\partial w_{11}} L(y, f(x)) \equiv \frac{\partial}{\partial w_{11}} (y - f(x))^2.$$

Ignore the case where $f(x)$ might not be differentiable (because the rectilinear function is not differentiable everywhere).

Hint: write out the formula using the chain rule and use the following definition of the derivative of $\sigma(s)$:

$$\frac{\partial}{\partial s} \sigma(s) = \sigma(s)(1 - \sigma(s)).$$

Solution. We first note that only the first hidden node h_1 depends on w_{11} , so we can write out the chain rule of derivatives from $L(y, f(x))$ to w_{11} through h_1 :

$$\frac{\partial}{\partial w_{11}} L(y, f(x)) = \frac{\partial L}{\partial f} \frac{\partial f}{\partial h_1} \frac{\partial h_1}{\partial w_{11}}. \quad (40)$$

We can expand each term in (40) as:

$$\begin{aligned}\frac{\partial L}{\partial f} &= -2(y - f), \\ \frac{\partial f}{\partial h_1} &= \frac{\partial \sigma(u_1 h_1 + u_2 h_2)}{\partial h_1} = \sigma(u_1 h_1 + u_2 h_2)(1 - \sigma(u_1 h_1 + u_2 h_2))u_1, \\ \frac{\partial h_1}{\partial w_{11}} &= \begin{cases} x_1 & \text{if } w_{11}x_1 + w_{21}x_2 > 0 \\ 0 & \text{otherwise} \end{cases}.\end{aligned}$$

(In the case where $w_{11}x_1 + w_{21}x_2 = 0$, then $\partial h_1 / \partial w_{11}$ is undefined because $\tau(0)$ is not differentiable. However, in practice this basically never happens, and so we generally ignore this case during stochastic gradient descent.)

8 Convolutional Kernels

(The purpose of this question is to improve your comfort in reasoning about spatial models such as convolutional networks.)

Here are three convolution filters:

$$K_1 = \begin{bmatrix} 0 & 0 & -0.375 & 0 & 0 \\ 0 & -0.375 & -0.75 & -0.375 & 0 \\ -0.375 & -0.75 & 7 & -0.75 & -0.375 \\ 0 & -0.375 & -0.75 & -0.375 & 0 \\ 0 & 0 & -0.375 & 0 & 0 \end{bmatrix},$$

$$K_2 = \begin{bmatrix} 0.0054 & 0.0180 & 0.0268 & 0.0180 & 0.0054 \\ 0.0180 & 0.0597 & 0.0890 & 0.0597 & 0.0180 \\ 0.0268 & 0.0890 & 0.1328 & 0.0890 & 0.0268 \\ 0.0180 & 0.0597 & 0.0890 & 0.0597 & 0.0180 \\ 0.0054 & 0.0180 & 0.0268 & 0.0180 & 0.0054 \end{bmatrix},$$

$$K_3 = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & -1 & 0 & 1 & 0 \\ 0 & -1 & 0 & 1 & 0 \\ 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}.$$

This question involves convolutional filters. Let K be a $(2r + 1) \times (2r + 1)$ -dimensional convolutional filter. Convolving an image X with K corresponds to a new image \tilde{X} with pixel values:

$$\tilde{X}_{i,j} = \sum_{i'=-r}^r \sum_{j'=-r}^r X_{i+i',j+j'} K_{i'+r+1,j'+r+1}.$$

In other words, each pixel $\tilde{X}_{i,j}$ in the convolved image is computed via the weighted sum of an image patch of X centered at $X_{i,j}$ with weights K (ignore boundary cases).



Figure 3: Original Image

Question: Suppose we convolve the three filters K_1 , K_2 , and K_3 on the image shown in Figure 3. Which resulting image in Figure 4 corresponds to which kernel?



Figure 4: Convolved Images

Solution. K_1 corresponds to the right image in Figure 4. K_1 is a sharpening filter, because it accentuates pixels that have high contrast with neighboring pixels.

K_2 corresponds to the middle image in Figure 4. K_2 is a blurring filter, because it computes a weighted average of a neighborhood of pixels.

K_3 corresponds to the left image in Figure 4. K_3 is an edge detection (or Gabor) filter, because it only activates when there is an increase in pixel intensity from left-to-right. Note that the sum of K_3 is 0, meaning that the resulting image \tilde{X} is black whenever there are no pixel intensity gradients in the image neighborhood patch.

9 Hidden Markov Models

(The purpose of this question is to improve your comfort in working with structured or graphical models.)

Given an input sequence $x = (x^1, \dots, x^M)$ and output sequence $y = (y^1, \dots, y^M)$, a hidden Markov model decomposes the probability of $P(x, y)$ as:

$$P(x, y) = \prod_{i=1}^M P(x^i | y^i) P(y^i | y^{i-1}), \quad (41)$$

where $y^0 = 0$ denotes a special start state.

Suppose each $y^i \in \{1, 2\}$ can take one of two states, and each $x^i \in \{A, B\}$ can take one of two observations. The $P(y^i | y^{i-1})$ probability table is:

$P(y^i y^{i-1})$	$y^i = 1$	$y^i = 2$
$y^{i-1} = 0$	2/3	1/3
$y^{i-1} = 1$	3/4	1/4
$y^{i-1} = 2$	1/2	1/2

and the $P(x^i | y^i)$ probability table is:

$P(x^i y^i)$	$x^i = A$	$x^i = B$
$y^i = 1$	2/3	1/3
$y^i = 2$	1/5	4/5

Question: Compute $P(x = (A, A, B))$ and $\operatorname{argmax}_y P(x = (A, A, B), y)$.

Solution. We can compute $P(x = (A, A, B))$ by marginalizing out all possible y 's:

$$P(x = (A, A, B)) = \sum_y P(x = (A, A, B), y), \quad (42)$$

and then we can just use (41).

OPTION 1: Brute Force Exhaustive Enumeration of All Possible y 's. This problem is small enough to exhaustively enumerate all possible y 's. There are 8 possible y 's that are length-3:

$$y \in \{(1, 1, 1), (1, 1, 2), (1, 2, 1), (1, 2, 2), (2, 1, 1), (2, 1, 2), (2, 2, 1), (2, 2, 2)\}.$$

One could simply brute force compute $P(x = (A, A, B))$ via (42) by summing over the 8 possible y 's:

$$\begin{aligned} P(x = (A, A, B), y = (1, 1, 1)) &= \frac{2}{3} \frac{2}{3} \frac{2}{3} \frac{1}{3} \frac{1}{3} = \frac{72}{1296}. \\ P(x = (A, A, B), y = (1, 1, 2)) &= \frac{2}{3} \frac{2}{3} \frac{2}{3} \frac{4}{5} \frac{1}{4} = \frac{96}{2160}. \\ P(x = (A, A, B), y = (1, 2, 1)) &= \frac{2}{3} \frac{2}{3} \frac{1}{5} \frac{1}{4} \frac{1}{3} = \frac{4}{1080}. \end{aligned}$$

$$P(x = (A, A, B), y = (1, 2, 2)) = \frac{2}{3} \frac{2}{3} \frac{1}{5} \frac{1}{4} \frac{1}{5} \frac{1}{2} = \frac{16}{1800}.$$

$$P(x = (A, A, B), y = (2, 1, 1)) = \frac{1}{5} \frac{1}{3} \frac{2}{3} \frac{1}{2} \frac{1}{3} \frac{1}{4} = \frac{6}{1080}.$$

$$P(x = (A, A, B), y = (2, 1, 2)) = \frac{1}{5} \frac{1}{3} \frac{2}{3} \frac{1}{2} \frac{1}{5} \frac{1}{4} = \frac{8}{1800}.$$

$$P(x = (A, A, B), y = (2, 2, 1)) = \frac{1}{5} \frac{1}{3} \frac{1}{5} \frac{1}{2} \frac{1}{3} \frac{1}{2} = \frac{1}{900}.$$

$$P(x = (A, A, B), y = (2, 2, 2)) = \frac{1}{5} \frac{1}{3} \frac{1}{5} \frac{1}{2} \frac{1}{5} \frac{1}{2} = \frac{4}{1500}.$$

The sum over the above eight values yields:

$$P(x = (A, A, B)) = \frac{72}{1296} + \frac{96}{2160} + \frac{4}{1080} + \frac{16}{1800} + \frac{6}{1080} + \frac{8}{1800} + \frac{1}{900} + \frac{4}{1500} \approx 0.126.$$

and taking the max over the above eight values yields:

$$\operatorname{argmax}_y P(x = (A, A, B), y) = (1, 1, 1).$$

OPTION 2a: Dynamic Programming via Forward Algorithm to Compute $P(x = (A, A, B))$. The dynamic programming solution leverages the following observation for length- M inputs x :

$$P(x) = \sum_y P(x, y) \tag{43}$$

$$= \sum_y \prod_{i=1}^M P(x^i | y^i) P(y^i | y^{i-1}) \tag{44}$$

$$= \sum_{y^{(1:M-1)}} \prod_{i=1}^{M-1} P(x^i | y^i) P(y^i | y^{i-1}) \sum_{y^M} P(x^M | y^M) P(y^M | y^{M-1}) \tag{45}$$

$$= \sum_{y^{(1:M-1)}} P(x^{(1:M-1)}, y^{(1:M-1)}) \sum_{y^M} P(x^M | y^M) P(y^M | y^{M-1}), \tag{46}$$

where $y^{(1:M-1)}$ denotes the length- $(M-1)$ prefix subsequence of y (i.e., the first $M-1$ tokens).

(46) defines a recursive definition of $P(x)$ that we can exploit to more efficiently and compactly compute $P(x)$. For a given x , define α_t^i as:

$$\alpha_t^i = \sum_{y^{(1:i-1)}} P(x^{(1:i)}, y^{(1:i-1)} \oplus t) \equiv P(x^{(1:i)} | y^i = t),$$

where $y^{(1:i-1)} \oplus t$ denotes the list append operator that appends token t to the end of $y^{(1:i-1)}$. In other words α_t^i is the total joint probability of the first i tokens in x , $x^{(1:i)}$, conditioned on $y^i = t$. Thus, we can write $P(x)$ for length- M x as:

$$P(x) = \sum_t P(x | y^M = t) = \sum_t \alpha_t^M. \tag{47}$$

Exploiting (46), we can recursively define α_t^i as:

$$\alpha_t^i = P(x^i | y^i = t) \sum_{t'} \alpha_{t'}^{i-1} P(y^i = t | y^{i-1} = t'), \quad (48)$$

which gives us an efficient and compact way to recursively compute each α^i vector and finally (47):

$$\begin{aligned} \alpha^1 &= \begin{bmatrix} \frac{2}{3} \frac{2}{3} \\ \frac{1}{5} \frac{1}{3} \end{bmatrix} \approx \begin{bmatrix} 0.4444 \\ 0.0667 \end{bmatrix}, \\ \alpha^2 &= \begin{bmatrix} \frac{2}{3} (\frac{3}{4} \alpha_1^1 + \frac{1}{2} \alpha_2^1) \\ \frac{1}{5} (\frac{1}{4} \alpha_1^1 + \frac{1}{2} \alpha_2^1) \end{bmatrix} \approx \begin{bmatrix} 0.2444 \\ 0.0289 \end{bmatrix}, \\ \alpha^3 &= \begin{bmatrix} \frac{1}{3} (\frac{3}{4} \alpha_1^2 + \frac{1}{2} \alpha_2^2) \\ \frac{4}{5} (\frac{1}{4} \alpha_1^2 + \frac{1}{2} \alpha_2^2) \end{bmatrix} \approx \begin{bmatrix} 0.0659 \\ 0.0604 \end{bmatrix}. \end{aligned}$$

Via (47) summing over α^3 yields $P(x = (A, A, B)) = \alpha_1^3 + \alpha_2^3 \approx 0.126$.

OPTION 2b: Dynamic Programming via Viterbi to Compute $\operatorname{argmax}_y P(x = (A, A, B), y)$. Define \hat{y}_t^i as the length- i solution ending in the token t that maximizes probability of $P(x^{(1:i)}, \hat{y}_t^i)$:

$$\hat{y}_t^i = \left(\operatorname{argmax}_{y^1, \dots, y^{i-1}} P(x^{(1:i)}, y^{(1:i-1)} \oplus t) \right) \oplus t. \quad (49)$$

We can also keep track of the probability as α_t^i :

$$\alpha_t^i = P(x^{(1:i)}, \hat{y}_t^i). \quad (50)$$

Assuming we have computed \hat{y}_t^M , the most likely y can be solved via:

$$\operatorname{argmax}_y P(x, y) = \operatorname{argmax}_t P(x, \hat{y}_t^M).$$

The key observation is that we can compute \hat{y} and α recursively via:

$$\hat{y}_t^i = \left(\operatorname{argmax}_{\hat{y}_{t'}^{i-1}} P(x^{(1:i)}, \hat{y}_{t'}^{i-1} \oplus t) \right) \oplus t, \quad (51)$$

$$\alpha_t^i = \operatorname{argmax}_{t'} \alpha_{t'}^{i-1} P(x^i | t) P(t | t'). \quad (52)$$

Typically (51) is computed as a side product of computing (52).

$$\begin{aligned} \alpha^1 &= \begin{bmatrix} \frac{2}{3} \frac{2}{3} \\ \frac{1}{5} \frac{1}{3} \end{bmatrix} \approx \begin{bmatrix} 0.444 \\ 0.067 \end{bmatrix}, \quad \hat{y}^1 = \begin{bmatrix} (1) \\ (2) \end{bmatrix}. \\ \alpha^2 &= \begin{bmatrix} \frac{2}{3} \frac{3}{4} \alpha_1^1 \\ \frac{1}{5} \frac{1}{4} \alpha_1^1 \end{bmatrix} = \begin{bmatrix} 0.222 \\ 0.022 \end{bmatrix}, \quad \hat{y}^2 = \begin{bmatrix} (1, 1) \\ (1, 2) \end{bmatrix}. \\ \alpha^3 &= \begin{bmatrix} \frac{1}{3} \frac{3}{4} \alpha_1^2 \\ \frac{4}{5} \frac{1}{4} \alpha_1^2 \end{bmatrix} = \begin{bmatrix} 0.056 \\ 0.044 \end{bmatrix}, \quad \hat{y}^3 = \begin{bmatrix} (1, 1, 1) \\ (1, 1, 2) \end{bmatrix}. \end{aligned}$$